


Detectors: Harris, Scale
invariance, Harris-Laplace,
affine Harris-Laplace

Example: Finding Correspondences Between Images



- First step toward 3-D reconstruction: Find correspondences between *feature points* in two images of a scene
- Object recognition: Find correspondences between *feature points* in “training” and “test” image





2,106 images, 819,242 points

<http://grail.cs.washington.edu/projects/rome/>

Building Rome in a Day


Sameer Agarwal, Noah Snavely, Ian Simon, Steven M. Seitz and Richard Szeliski

International Conference on Computer Vision, 2009, Kyoto, Japan.




Interest points as “corner-like”:
Basic Harris detector




Interest Points






- Intuitively, junctions of contours.
- Generally more stable features over changes of viewpoint
- Intuitively, large variations in the neighborhood of the point in all directions



The distribution of the x and y derivatives is very different for all three types of patches



The distribution of x and y derivatives can be characterized by the shape and size of the principal component ellipse



How to evaluate “interestness”?

- The distribution of gradients in a neighborhood W is represented by the inertia (shape, Harris) matrix:

$$H = \begin{bmatrix} \sum_W I_x^2 & \sum_W I_x I_y \\ \sum_W I_x I_y & \sum_W I_y^2 \end{bmatrix}$$

- Elongations of the distribution = Eigenvalues of H : λ_{\min} , λ_{\max}
 - We want λ_{\min} , λ_{\max} to be approx. equal
 - We want λ_{\min} , λ_{\max} to be large

Harris detector and its variants

- We want λ_{\min} , λ_{\max} to be approx. equal
- We want λ_{\min} , λ_{\max} to be large

$$\mathbf{R} = 4 \frac{\lambda_{\min} \lambda_{\max}}{(\lambda_{\min} + \lambda_{\max})^2}$$

- ($R = 1$ if $\lambda_{\min} = \lambda_{\max}$ but keep only the ones with large λ_{\max})

$$\mathbf{R} = \frac{\lambda_{\min} \lambda_{\max}}{(\lambda_{\min} + \lambda_{\max})}$$

$$\mathbf{R} = \lambda_{\min} \lambda_{\max} - k(\lambda_{\min} + \lambda_{\max})^2$$

Comp. efficient definition

- For any symmetric matrix H :
 - $\text{Det}(H) = \lambda_{\min} \lambda_{\max}$
 - $\text{Trace}(H) = \lambda_{\min} + \lambda_{\max}$ (The trace is the sum of the diagonal of H)

$$R = 4 \frac{\text{Det}(H)}{\text{Trace}(H)^2}$$

$$R = \frac{\text{Det}(H)}{\text{Trace}(H)}$$


$$R = \text{Det}(H) - k \text{Trace}(H)^2$$

Computation of H and gradients

- I_x means convolution with Gaussian of σ
- H should be computed with different weights \rightarrow Convolution with Gaussian

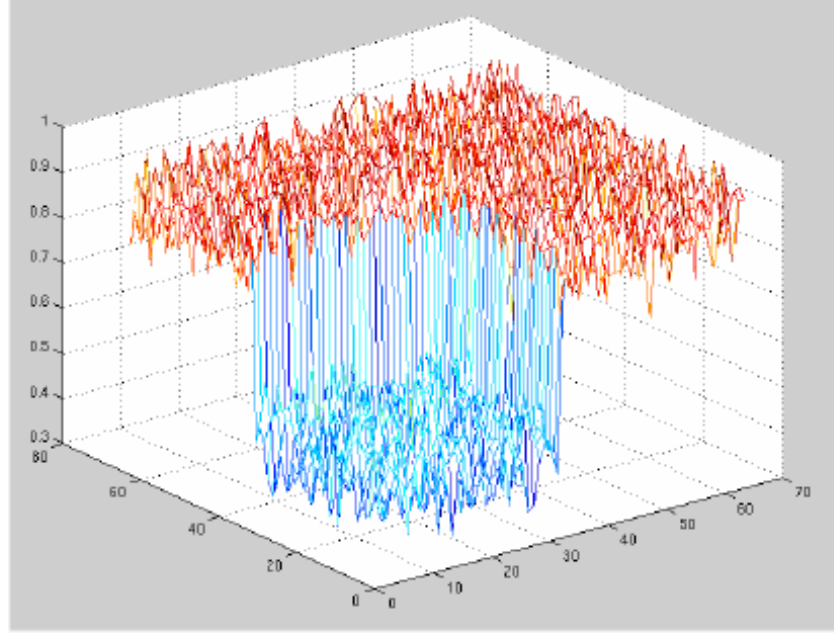



$$H = \begin{bmatrix} \sum_W I_x^2 & \sum_W I_x I_y \\ \sum_W I_x I_y & \sum_W I_y^2 \end{bmatrix}$$




$$H = \begin{bmatrix} G_\sigma * I_x^2 & G_\sigma * I_x I_y \\ G_\sigma * I_x I_y & G_\sigma * I_y^2 \end{bmatrix} = G_\sigma * \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$



Input




Response of interest
point operator



Input



Response of interest
point operator



Interest point detection example

Interest points example



1. Compute x and y derivatives of image

$$I_x = G_{\sigma}^x * I \quad I_y = G_{\sigma}^y * I$$

2. Compute products of derivatives at every pixel

$$I_{x2} = I_x \cdot I_x \quad I_{y2} = I_y \cdot I_y \quad I_{xy} = I_x \cdot I_y$$

3. Compute the sums of the products of derivatives at each pixel

$$S_{x2} = G_{\sigma^2} * I_{x2} \quad S_{y2} = G_{\sigma^2} * I_{y2} \quad S_{xy} = G_{\sigma^2} * I_{xy}$$

4. Define at each pixel (x, y) the matrix



$$H(x, y) = \begin{bmatrix} S_{x2}(x, y) & S_{xy}(x, y) \\ S_{xy}(x, y) & S_{y2}(x, y) \end{bmatrix}$$

5. Compute the response of the detector at each pixel

$$R = \text{Det}(H) - k(\text{Trace}(H))^2$$

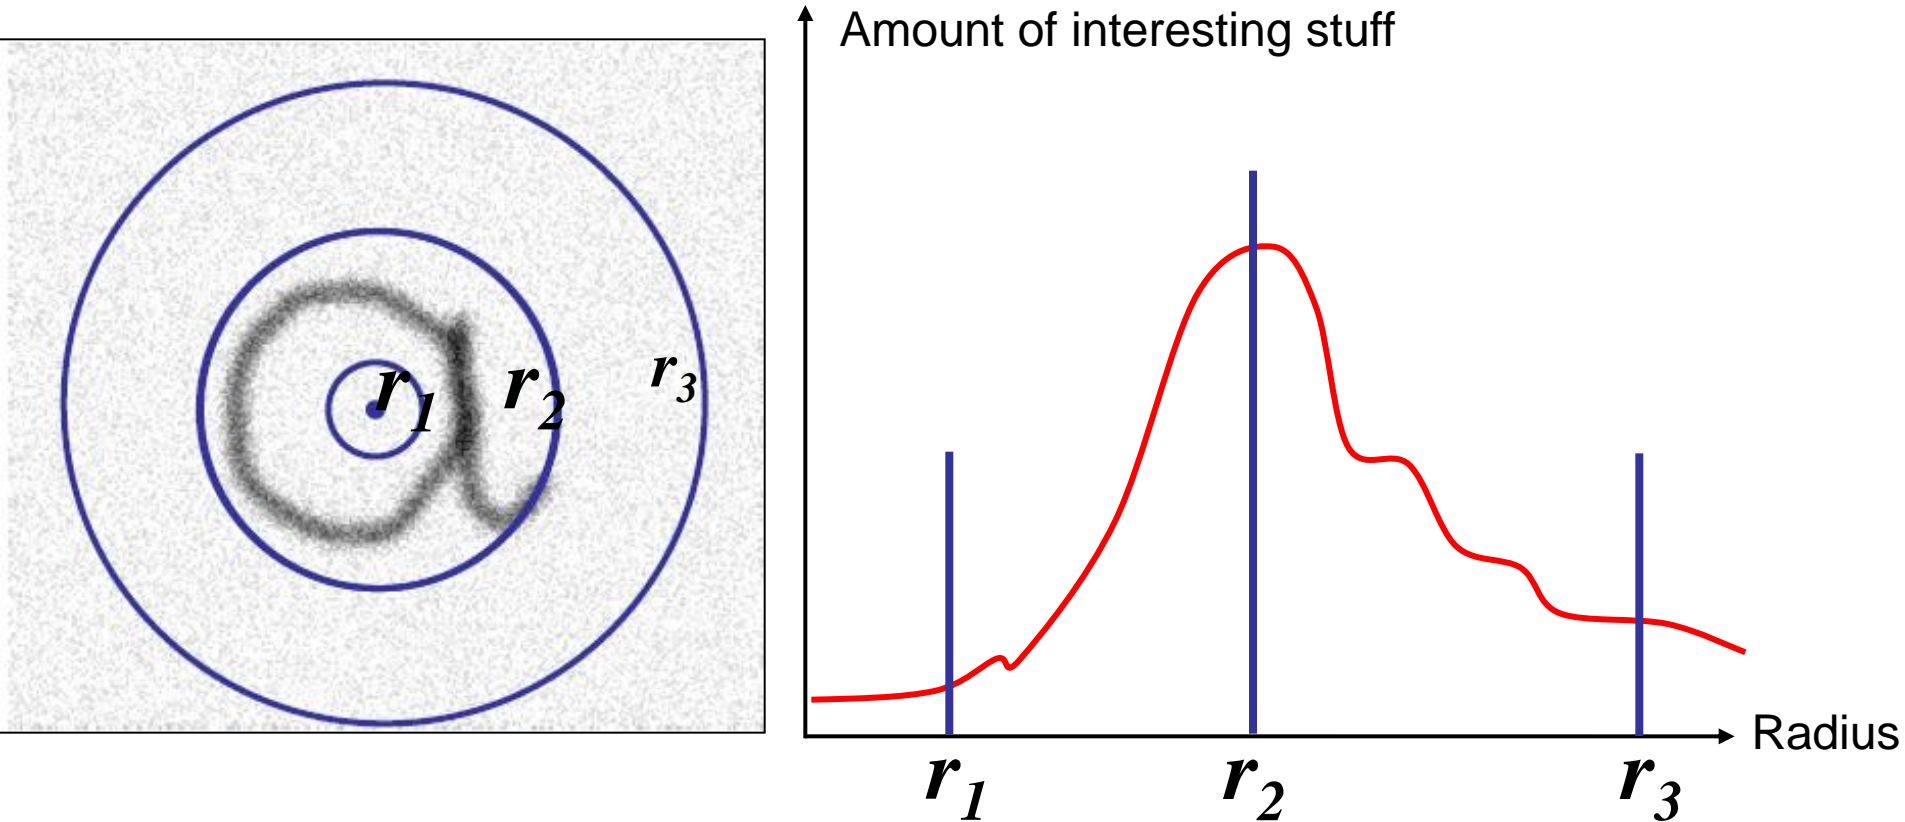
Characteristic scale

Scale Selection



- We need to decide what window size to use for computing the Harris matrix
- Equivalently, we need to choose the value of σ'
- The window size (or σ') must be consistent between different magnifications of the image

How to choose a neighborhood size? → Intuition





Best radius: Local extrema of function that measures amount of interesting stuff



X 2

$$\sigma_1 = 2 \times \sigma_2 !!!$$



- Why did we use the normalized Laplacian in the previous example?
- Justification (and basis for most scale selection operations in computer vision):
- *Scale Selection Principle (T. Lindeberg):*

In the absence of other evidence, assume that a scale level, at which some (possibly non-linear) combination of normalized derivatives assumes a local maximum over scales, can be treated as reflecting a characteristic length of a corresponding structure in the data.

What are normalized derivatives?

$$\sigma^{n+m} \frac{\partial^{n+m} f}{\partial x^n \partial y^m}$$

Example using 2nd order derivatives

$$\sigma^2 \nabla^2 f = \sigma^2 \left(\frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2} \right)$$

$$\nabla_{\sigma}^2 I = \nabla^2 G_{\sigma} * I$$

Constant independent of σ






$$G_{k\sigma} - G_{\sigma} \approx (k-1)\sigma^2 \nabla^2 G_{\sigma}$$

The Laplacian of a Gaussian can be approximated by the difference of two Gaussians.

To compare the Laplacian at different scales we need to explain more carefully what the approximation is.




In practice: the scaled Laplacian can be computed by taking the difference between level in a Gaussian pyramid





(From Mikolajczyk
and Schmid'02)






Laplacian

Interest points as “blob-like”: Basic
Laplacian detector = Local maxima
of characteristic scale

From Laplacian pyramid to detected points

- Find local extrema of $\nabla^2 G_\sigma$ over x, y , and σ (scale)



Combining the two: Harris-
Laplace detector

Combining Harris and Laplace:

Harris-Laplace detector

$$H = G_{\sigma_I} * \begin{bmatrix} I_{x,\sigma_D}^2 & I_{x,\sigma_D} I_{y,\sigma_D} \\ I_{x,\sigma_D} I_{y,\sigma_D} & I_{y,\sigma_D}^2 \end{bmatrix} \quad R = \text{Det}(H) - k(\text{Trace}(H))^2$$

- The location x of the points detected by Harris is not scale invariant \rightarrow Depends on the choice of σ_I and σ_D .
- Reduce to one parameter: $\sigma_D = s\sigma_I$ ($s = 0.7$)
- The Laplacian trick gives us a good σ but not where the interest point is.
- Chicken and egg problem:
 - If we knew x we could estimate σ
 - If we knew σ we could find x

Harris-Laplace: Algorithm *summary*


1. Try Harris at different scales and report initial points and associated scales



- Try different σ_l of the form $k^n \sigma_o$ ($k = \sqrt{2}$)
- Report points with large R

2. For each detected point x

- Estimate characteristic scale σ_c as maximum of $\sigma^2 \nabla^2 \mathbf{G}_\sigma$
- Find the point x' with the maximum of R in a 8x8 neighborhood of x by using the new scale $\sigma_l = \sigma_c$
- Replace x by x'

Iterate until convergence









Example from *Mikolajczyk and Schmid 2004*

Affine-invariant detection (*overview only*)



- Need to define a richer description of “neighborhood” or scale
- Use directional derivatives instead: σ_I replaced by Σ_I , σ_D replaced by Σ_D
- Σ_I represent an elliptical “neighborhood” instead of a circular one
- More degrees of freedom to search through but conceptually similar algorithm:
- Assume $\Sigma_D = s\Sigma_I$
- Find x ’s with initial selections of Σ_I
- Iterate:
 - Re-estimate “scale” Σ_I
 - Adjust the location of x based on new Σ_I

State of the Art: Affine Invariance




Example from *Mikolajczyk and Schmid 2004*

Affine Invariance



Application: Finding Correspondences





Initial detections

Scale: 4.9

Rotation: 19°

Example from *Mikolajczyk and Schmid 2004*




Final matches: 32 correct correspondences

Scale: 4.9

Rotation: 19°

Example from *Mikolajczyk and Schmid 2004*

Application: Finding Correspondences



Scale change: 1.7

Viewpoint change: 50°

Example from *Mikolajczyk and Schmid 2004*